



Research Supervisor:
Implementation Supervisor:
Department of Computer Science · Faculty of Engineering and Science ·
May 2025

Problem & Motivation


Language-learning platforms rely on **pre-simplified materials**. Authentic English — films, podcasts, news — is linguistically rich but cognitively inaccessible without structural guidance. Simplification destroys the input; learners never encounter real language.

Hypothesis: Deterministic NLP and lightweight generative AI can *explain* authentic input without altering it — producing hierarchical annotations that scaffold meaning while preserving full linguistic complexity.

- How can a hybrid NLP + LLM pipeline produce accurate pedagogical annotations from authentic media?
- What data model captures recursive linguistic relations while remaining human-interpretable?
- Can bilingual visualisation reduce cognitive load?


The RLE Data Model

Core innovation: every sentence is a **nested JSON hierarchy** at three levels — Sentence → Phrase → Word — validated against a dual-schema contract before persistence.




Schema is **language-agnostic**. Current focus: English → Russian. Pluggable translator backends support additional target languages.

Analysis Pipeline



Research Journey: What We Tried


Each component went through multiple candidate approaches before the final solution was selected:



The New Physics of Context-Pattern Interaction

"In open systems, context forms patterns. In closed systems, patterns form context."
— Principle discovered empirically during T5 fine-tuning

The Duality Loop Diagram



The Key Insight: Placeholder-Driven Learning

Fine-tuning T5 on raw linguistic notation **failed to generalise**. The breakthrough was replacing lexical tokens with structural placeholders:

✗ Raw: "She gave him a book" → [VP gave][NP him]
✓ Template: [SUBJ][V-PAST][OBJ][OBJ1][NOUN]
→ [TAM: past-3sing-he][VOICE: active]

The abstract grammatical **pattern** — not the words — became the unit of learning. The model generalises to novel vocabulary because the pattern is *invariant to lexical substitution*.

This is closed-system behaviour: the internal pattern imposes structure on any incoming context — regardless of whether that context was seen during training.


Implication for Agentic AI

An LLM agent navigating the real world is a **closed system** operating in an **open environment**. Its trained patterns are fixed; the world's context is not. The duality principle predicts exactly where agents succeed (stable structural patterns) and fail (distribution shift beyond the trained pattern's boundary).

Understanding training as pattern-carving and inference as pattern-imposing offers a new lens for agent reliability research.

Training Data & Corpus Scale


Four open-access corpora merged into a CEFR-labelled full-ladder dataset (A1 → C2):



27,000 raw samples ingested
685 unique after deduplication
6 CEFR levels covered (A1 → C2)
Full-ladder per grammar class

Live Demo
www.el-a.uk
Try it in your browser

Performance Metrics



Annotation Backend Comparison

Approach	Quality	Cost	Offline
GPT-4o	High	\$\$\$	✗
Rules only	Medium	Free	✓
T5 fine-tuned	High	Free	✓
XGBoost CEFR	78.3%	Free	✓

T5 fine-tuning achieves GPT-quality annotation at zero inference cost, running entirely on-device.

System Validation

- O1 (RLE Generation):** spaCy + TAM engine produces schema-valid RLE for all test sentences ✓
- O2 (ML Enrichment):** CEFR 78.3%; T5 ROUGE-L 0.81 ✓
- O3 (UI & Export):** Linguistic Visualizer renders full RLE tree; JSON/CSV export verified ✓
- O4 (Deployment):** Android PWA install — fully client-side, no cloud dependency ✓

Conclusions

- Explainable AI annotation makes authentic English **pedagogically accessible without simplification**.
- The RLE schema is a viable, human-interpretable model of recursive linguistic structure.
- The **Context-Pattern Duality**, observed empirically during T5 fine-tuning, generalises beyond NLP to any learning system that transitions from open-world training to closed-system inference.
- Placeholder-driven T5 fine-tuning generalises to novel vocabulary because the abstract structural pattern is lexically invariant.

Future Directions

- Tauri 2 Desktop App:** Native offline package (Linux .deb / Windows) with bundled ML models — no browser required, full hardware access for on-device inference.
- Translation Quality Gates:** Round-trip Chrf score, length-ratio guards, n-gram repetition penalty, and a human review queue wired into the M2M100 pipeline.
- M2M100 LoRA Fine-Tuning:** Domain-specific refinement of the translation model on accepted English → Russian pairs from the review queue.
- Additional target languages:** Pluggable translator backends support any language beyond the current English → Russian focus.
- Real-time streaming ASR:** Live audio analysis for conversation practice via Whisper streaming mode.

References

1. Honnibal, M. & Montani, L. (2017). spaCy: Industrial-strength NLP.
2. Raffel, C. et al. (2020). Exploring the Limits of Transfer Learning with T5. JMLR.
3. Radford, A. et al. (2023). Robust Speech Recognition via Large-Scale Weak Supervision. ICML.
4. Chen, T. & Guestrin, C. (2016). XGBoost: A Scalable Tree Boosting System. ACM KDD.
5. Levy, M. (1997). *Computer-Assisted Language Learning*. Oxford University Press.

Acknowledgements

The author sincerely thanks **Prof. Nasir Ahmad** (Implementation Supervisor) and **Dr. Alex Vahabzadeh** (Research Supervisor) for their invaluable guidance throughout this project. Thanks to the Department of Computer Science at MTU Cork for the academic environment that made this work possible. Inspired by *Eleonora Rastvorova*.